

Autoregressive Generation of Neural Field Weights

Using a transformer based architecture

from Luis Muschal and Luca Fanelau

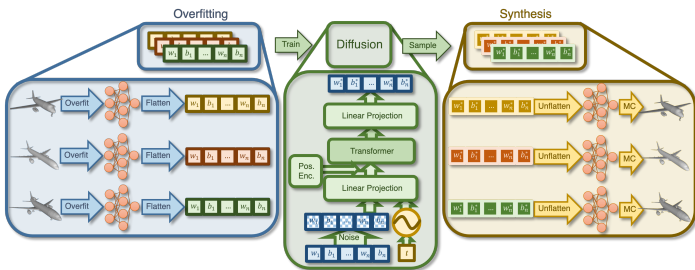
Table of Content

1. Related works
2. Autoregressive Generation of Neural Field Weights
3. Regression Transformer
4. Challenges: Permutation Symmetries
5. Overfitting Neural Fields
6. Regression Transformer
7. Conclusion and Outlook: Tokenization

Related works

HyperDiffusion

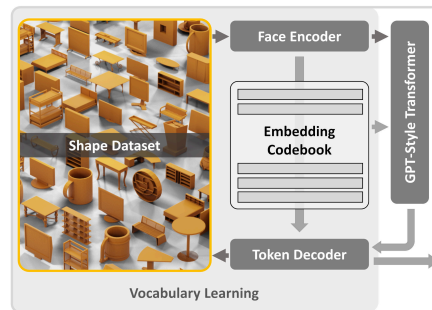
Operates on MLP weights directly to generate new neural implicit fields encoded by synthesized MLP parameters



ICCV'23 [Erkoç et al.]: Hyperdiffusion

MeshGPT

Autoregressively generate triangle meshes as sequences of triangles using a learned vocabulary of latent quantized embedding as tokens



CVPR'24 [Siddiqui et al.]: MeshGPT

Autoregressive Generation of Neural Field Weights

Neural Fields

Input coordinate location in n-dimensional space are mapped to target signal domain

Example:

With S being a surface in a 3-dimensional space \mathbb{R}^3 . The Signed Distance Function $f : \mathbb{R}^3 \rightarrow \mathbb{R}$ is defined for a point $\mathbf{p} \in \mathbb{R}^3$ as:


$$f_{\Theta}(\mathbf{p}) = \begin{cases} \text{distance}(\mathbf{p}, S) & \text{if } \mathbf{p} \text{ is outside } S, \\ 0 & \text{if } \mathbf{p} \text{ is on } S, \\ -\text{distance}(\mathbf{p}, S) & \text{if } \mathbf{p} \text{ is inside } S, \end{cases}$$



Autoregressive Generation of Neural Field Weights

Autoregressive Process

- Goal: generative modeling of neural fields $P(\theta_i \mid \theta_{i-1}, \theta_{i-2}, \dots, \theta_0)$
- Using a generally available preset for GPT-like Architecture (like nanoGPT)
- Use Transformer to sample from the Probability, eg. $\theta_i = \text{Transformer}(\theta_{i-1}, \theta_{i-2}, \dots, \theta_0)$



Autoregressive Generation of Neural Field Weights

From nanoGPT to Regression Transformer

nanoGPT

vs. Our Regression Transformer

Tokenizer

MLP Embedding on weight

Embedding + Positional Embedding


N x Blocks (Causal Self Attention and MLP)

Linear Transformation Embedding


Cross-Entropy Loss

L1-norm as Loss

Ground Truth



N=1




Iteration: 0

Regression Transformer

Observing the Effects of Increasing N

- Transformer fails to capture the structure of the weights for larger N
- Why can't the sequence be remembered even for small values of N?

Ground
Truth



N=4



N=32



Iteration: 0


Challenges: Permutation Symmetries

The same signal can be represented by different weight matrices

$$P = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

permuted weight matrices are calculated using:

$$\begin{aligned}\tilde{W}_0 &= PW_0 \\ \tilde{W}_1 &= W_1 P^T\end{aligned}$$



Overfitting Neural Fields

Finding a Solution

Minimize structural change by
conditioning the training process
using weight initialization

Approach:

Overfit single sample

Use weights for different sample
(conditioned)

Train sample on randomly initialized
weights (unconditioned)

Overfitting Neural Fields

Overfitting on one sample

Minimize structural change by
conditioning the training process
using weight initialization


Approach:

Overfit single sample


Use weights for different sample
(conditioned)

Train sample on randomly initialized
weights (unconditioned)


Ground Truth




First Sample



Legend




w_0



b_0



w_1




b_1



w_2



b_2



Epoch: 0

Overfitting Neural Fields

Overfitting on other sample

Minimize structural change by
conditioning the training process
using weight initialization


Approach:

Overfit single sample


Use weights for different sample
(conditioned)

Train sample on randomly initialized
weights (unconditioned)


Ground Truth






Unconditioned



Conditioned




Condition



Epoch: 0

Overfitting Neural Fields


Visualizing the Difference:




$$\Delta(W) = W_{\text{pretrained}} - W$$

$$\Delta(b) = b_{\text{pretrained}} - b$$


Ground Truth





Unconditioned



Conditioned



Condition




Epoch: 0

Reminder: Regression Transformer

How far we got with unconditioned neural fields

- Transformer fails to capture the structure of the weights for larger N
- Why can't the sequence be remembered even for small values of N?

Ground
Truth



N=4



N=32



Iteration: 0

Regression Transformer

Using conditioned neural fields to verify the Hypothesis

- Training Regression Transformer using conditioned Neural Fields Weights
- Structural similarity of weights improve the performance of the Transformer

Ground
Truth



N=4



N=32



Iteration: 0

Conclusion and Outlook: Tokenization

Predicting the next MLP weight as a token

Run into issues regarding special tokens:

$$\theta_i = \text{Transformer}(\theta_{i-1}, \theta_{i-2}, \dots, \theta_0) \rightarrow \theta_0?$$

Solution: Find Tokens to encode the MLP weights and transfer from Regression Transformer to Classical Transformer Architectures

First Approach:

- Create Tokens using conditioned Neural Field Weights
- Naive Attempt: Use weight distribution for discretization
- Vector Quantization Attempt: Find optimal token representation using optimization techniques

Second Approach:

- Find layer representations of unconditioned neural fields that are permutation equivariant
- For example by using the graph structure of the neural fields and employing deep learning techniques suited for graphs

Thank you for your attention!

We hope you enjoyed our presentation and are looking forward to your questions.